

## **MEMORANDUM - Built but Broken: The Case for Coordinating AI Model Card Transparency Legislation**

**To:** Justin Bullock, Americans For Responsible Innovation

**From:** Vashishth Doshi, Saumya Goyal, Ziqian Zhong

**Date:** April 30, 2026

---

### **Executive Summary**

Model card transparency is not absent from law. It is present three times over, in three structurally distinct forms, moving toward the same destination without coordination. Across federal procurement rules, state-level algorithmic accountability statutes, and voluntary framework-based safe harbours, governments and regulators have independently arrived at the same underlying requirement — that AI developers disclose what their models are, how they were built, and what they can and cannot do. What they have not arrived at is a shared mechanism, a shared vocabulary, or a shared enforcement baseline. This memo maps that legislative architecture across three tiers, identifies the structural failure at each, and describes what connecting them would require.

### **Legislation with Implicit Incentives**

The broadest tier of model card-adjacent regulation operates through incentive rather than obligation. The NIST AI Risk Management Framework, adopted January 2023, recommends model cards within its Map and Measure functions as a tool for documenting intended use, known limitations, and performance characteristics across deployment contexts.<sup>1</sup> It carries no enforcement mechanism. Texas's Responsible AI Governance Act (TRAIGA, 2025) translates this into a legal structure by designating NIST RMF compliance as an affirmative defence against liability — meaning developers who produce NIST-aligned documentation, including model cards, are partially shielded from litigation.<sup>2</sup> The incentive is real, but its activation depends entirely on whether a developer calculates that litigation risk exceeds compliance cost. For large, well-resourced frontier model developers, that calculation has historically favoured non-disclosure: the litigation exposure under TRAIGA remains untested, enforcement operationalization is at zero, and the reputational cost of absent documentation has not yet materialised as a financial one. Voluntary frameworks, by design, protect willing adopters, but don't move unwilling ones.

### **Legislation with Disclosure Mandates**

The middle tier is where transparency becomes enforceable. A growing cluster of US state statutes and international instruments require developers and deployers to publish specific, enumerable disclosures about their AI systems — without naming model cards as the vehicle. Colorado's AI Act (SB 24-205, effective June 2026) requires impact assessments documenting algorithmic discrimination risk for any high-risk AI system touching employment, housing, credit, or education.<sup>3</sup> New York City's Local Law 144 (in force since July 2023) requires independent bias audits with publicly disclosed results disaggregated by race, ethnicity, and sex for any automated employment decision tool.<sup>4</sup> Illinois's Human Rights Act amendment (HB 3773, effective January 2026) creates civil rights liability for discriminatory AI employment decisions, making the absence of bias documentation legally material in litigation.<sup>5</sup> At the international level, the EU AI Act's high-risk provisions under Annex III impose analogous documentation obligations across credit, education, employment, and critical infrastructure.<sup>6</sup> Laid side by side, these instruments require disclosure of training data provenance, performance disaggregation,

known limitations, and intended use — the four core sections of a standard model card.<sup>7</sup> A developer operating across all four jurisdictions simultaneously is meeting non-identical versions of the same underlying requirement, with no shared template, no shared vocabulary, and no mechanism for one jurisdiction's compliance to satisfy another's. The result is a paradox: the content mandate tier generates more documentation than any other, and less legible transparency than any other.

### **Legislation with Explicit 'Model Card' Mandates**

The uppermost tier is the only one where model card-equivalent documentation is required by name or direct structural specification, with minimum content defined and enforceable consequence attached. Three instruments currently occupy this rung. OMB Memorandum M-26-04 (December 2025) designates Model, System, or Data Cards as minimum transparency deliverables in all federal large language model procurement contracts, making their absence grounds for contract termination.<sup>8</sup> The EU's General-Purpose AI Code of Practice (August 2025), issued under Article 56 of the EU AI Act, adopts the standard model card template as the reference compliance format for all GPAI providers, with adherence creating a legal presumption of AI Act conformity.<sup>9</sup> California's Transparency in Frontier AI Act (SB 53, effective September 2025) requires frontier developers — those training models above  $10^{26}$  floating point operations — to publish structured safety frameworks covering catastrophic risk prevention, incident reporting, and capability limitations.<sup>10</sup> SB 53's threshold mirrors the EU AI Act's systemic risk definition precisely, and its disclosure obligations map directly to the safety and limitations sections of a standard model card. Together these three instruments represent the only regulatory contexts in which a developer faces a specified documentation obligation rather than a generalised disclosure duty.

The coverage gap, however, is significant. OMB M-26-04 reaches only vendors contracting directly with US federal agencies.<sup>11</sup> The GPAI Code of Practice reaches only providers above defined thresholds accessing the EU market. SB 53, while broader than either, applies only to the narrowest band of frontier developers by compute threshold, and its enforcement mechanisms remain untested. The SOTA consumer-facing models most consequential to public life — those audited in the accompanying SOTA model card audit-analysis — sit at the edge of all three perimeters simultaneously. The most visible AI systems in deployment face the most precisely written, and the most sparsely applied, documentation requirements currently in force.

### **The Gaps in this Legislation Ladder**

Taken together, the three tiers describe a regulatory architecture that is simultaneously over-built and under-connected. Each rung addresses the enforcement failure of the rung below it — content mandates make implicit incentives binding, explicit mandates make content requirements legible — but introduces a new structural failure of its own. The implicit tier produces no documentation. The content mandate tier produces documentation that cannot be compared across jurisdictions. The explicit mandate tier produces comparable, structured documentation that reaches only a narrow slice of the models actually in deployment. The net result is that the developers most capable of absorbing compliance costs — large, well-resourced frontier model labs — face the weakest aggregate documentation pressure, while smaller deployers operating in high-risk domains face arguably the heaviest, most fragmented, and least standardised set of disclosure obligations. What the ladder reveals, in other words, is not a progression toward transparency but a seemingly unfair distribution of regulatory burden that is inversely related to the scale of public consequence.

### **Advocating for Best Practices – for Americans for Responsible Innovation**

First, explicit naming. Content mandate jurisdictions require the substance of a model card without requiring the artefact. The result is that developers produce fragmented, jurisdiction-specific disclosures that satisfy the letter of each statute and the spirit of none. OMB M-26-04 demonstrates that naming the artefact directly — designating Model, System, or Data Cards as a defined deliverable with specified minimum content — resolves this ambiguity at a stroke. Extending that naming convention into state-level content mandate statutes, by amendment or by regulatory guidance, would immediately convert existing disclosure obligations into a coherent documentation regime.

Second, a standardised template built into legislation. The GPAI Code of Practice is the only quasi-legislative instrument that moves beyond naming to structure — it reproduces a defined set of fields, derived from the original model card specification, that providers must populate. That approach should inform binding legislation. A template adopted by reference across jurisdictions would collapse the fragmentation problem the disclosure-content mandate tier currently creates.

Third, minimum qualitative depth as a concurrent best-practice standard. Explicit naming and template standardisation address form. They do not address substance. The accompanying audit of six state-of-the-art models finds that even where model cards exist, they routinely omit or underspecify training data provenance, disaggregated performance evaluation, and known limitations — the fields most consequential for downstream accountability (see Part II: SOTA Model Card Audit, Zhong Z., Doshi, V., Goyal, S., 2026). Legislative instruments currently specify that these fields must be present. None specifies what adequate population of those fields looks like. ARI is positioned to advance a minimum qualitative depth standard as a parallel best-practice effort, developed in coordination with developers, civil society, and standards bodies, that gives content to the form that legislation is beginning to require.

---

<sup>1</sup> National Institute of Standards and Technology. (2023). Artificial intelligence risk management framework (AI RMF 1.0). U.S. Department of Commerce. <https://airc.nist.gov/airmf-resources/airmf/>

<sup>2</sup> Texas Legislature. (2025). House Bill 149: Texas Responsible AI Governance Act (TRAIGA). 89th Texas Legislature. <https://capitol.texas.gov/BillLookup/History.aspx?LegSess=89R&Bill=HB149>

<sup>3</sup> Colorado General Assembly. (2024). Senate Bill 24-205: Colorado Artificial Intelligence Act. <https://leg.colorado.gov/bills/sb24-205>

<sup>4</sup> New York City Council. (2021). Local Law 144: Automated employment decision tools. NYC Legistar. <https://legistar.council.nyc.gov/LegislationDetail.aspx?GUID=B051915D-A9AC-451E-81F8-6596032FA3F9&ID=4344524>

<sup>5</sup> Illinois General Assembly. (2024). House Bill 3773 / Public Act 103-0804. Illinois Legislative Information System. <https://www.ilga.gov/legislation/BillStatus.asp?DocNum=3773&GAID=17&DocTypeID=HB&SessionID=112&GA=103>

<sup>6</sup> European Parliament and Council of the European Union. (2024). Regulation (EU) 2024/1689 of the European Parliament and of the Council (EU AI Act), Annex III. Official Journal of the European Union. <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:32024R1689>

<sup>7</sup> Mitchell, M., Wu, S., Zaldivar, A., Barnes, P., Vasserman, L., Hutchinson, B., Spitzer, E., Raji, I. D., & Gebru, T. (2019). Model cards for model reporting. Proceedings of the 2019 ACM Conference on Fairness, Accountability, and Transparency (FAccT '19). <https://arxiv.org/abs/1810.03993>

<sup>8</sup> Office of Management and Budget. (2025, December 11). Memorandum M-26-04: Increasing public trust in artificial intelligence through unbiased AI principles. Executive Office of the President. <https://www.whitehouse.gov/wp-content/uploads/2025/12/M-26-04-Increasing-Public-Trust-in-Artificial-Intelligence-Through-Unbiased-AI-Principles-1.pdf>

<sup>9</sup> European Commission. (2025, July 10). General-purpose AI code of practice. European AI Office. <https://digital-strategy.ec.europa.eu/en/policies/contents-code-gpai>

<sup>10</sup> California Legislature. (2025). Senate Bill 53: Transparency in Frontier AI Act. California Legislative Information. [https://leginfo.ca.gov/faces/billNavClient.xhtml?bill\\_id=202520260SB53](https://leginfo.ca.gov/faces/billNavClient.xhtml?bill_id=202520260SB53)

<sup>11</sup> Office of Management and Budget. (2025, February). Memorandum M-25-22: Driving efficient acquisition of artificial intelligence in government. Executive Office of the President. <https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-22-Driving-Efficient-Acquisition-of-Artificial-Intelligence-in-Government.pdf>